


✓ Loading Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from transformers import AutoTokenizer, AutoModel, AutoModelForSequenceClassification
from torch.utils.data import DataLoader, Dataset
import torch
import torch.nn as nn
from datasets import Dataset, DatasetDict
```



✓ Loading Data

```
data_review = pd.read_json("yelp_academic_dataset_review.json", nrows=1000, lines=True)
```

```
data_review.head()
```

		review_id	user_id	business_id	stars
0		KU_O5udG6zpxOg-VcAEodg	mh_-eMZ6K5RLWhZylSBhwA	XQfwVwDr-v0ZS3_CbbE5Xw	
1		BiTunyQ73aT9WBnpR9DZGw	OyoGAe7OKpv6SyGZT5g77Q	7ATYjTlgM3jUlt4UM3lypQ	
2		saUsX_uimxRICVr67Z4Jig	8g_iMtfSiwikVnbP2etR0A	YjUWPpl6HXG530lwP-fb2A	
3		AqPFMleE6RsU23_auESxiA	_7bHUi9Uuf5__HHc_Q8guQ	kxX2SOes4o-D3ZQBkiMRfA	
4		Sx8TMOWLNuJBWer-0pcmoA	bcjbaE6dDog4jkNY91ncLQ	e4Vwtrqf-wpJfwesgvdgxQ	

```
data_review.shape
```

```
(50000, 9)
```

```
data_review.columns
```

```
Index(['review_id', 'user_id', 'business_id', 'stars', 'useful', 'funny',  
      'cool', 'text', 'date'],  
      dtype='object')
```

```
data = data_review.copy(deep=True)
```

✓ Preprocess Data

```
print("\nPreprocessing Text Data...")
data['text_cleaned'] = (
    data['text']
    .str.lower()
    .str.replace(r'http\S+', '', regex=True) # Remove URLs
    .str.replace(r'^\w\s', '', regex=True) # Remove punctuation
    .str.replace(r'\d+', '', regex=True) # Remove numbers
)
```



Preprocessing Text Data...

```
nltk.download('vader_lexicon')
```



```
[nltk_data] Downloading package vader_lexicon to
[nltk_data] /home/dgilkey/nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!
True
```

✓ Sentiment Mapping

```
def map_sentiment(stars):
    if stars > 2:
        return 2 # Positive
    elif stars == 2:
        return 1 # Neutral
    else:
        return 0 # Negative
```

```
data['sentiment_label'] = data['stars'].apply(map_sentiment)
```

```
data.sentiment_label.value_counts()
```



```
2    40618
0     5379
1     4003
Name: sentiment_label, dtype: int64
```

```
data.stars.value_counts()
```

```
↩ 5      22220
   4      12721
   3       5677
   1       5379
   2       4003
   Name: stars, dtype: int64
```

✓ Data Preparation

```
X = data['text_cleaned']
y = data['sentiment_label']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random
```

✓ Loading Transformer

```
print("\nLoading Transformer Model...")
model_name = "distilbert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModel.from_pretrained(model_name)
↩ /home/dgilkey/.local/lib/python3.10/site-packages/huggingface_hub/file_down
  warnings.warn(
```

```
    Loading Transformer Model...
```

```
def generate_embeddings(texts, tokenizer, model, max_length=128):
    """Generate sentence embeddings using a transformer model."""
    tokens = tokenizer(texts, padding=True, truncation=True, max_length=max_len
    with torch.no_grad():
        outputs = model(**tokens)

    embeddings = outputs.last_hidden_state[:, 0, :].numpy()
    return embeddings
```

✓ Embedding Generation

```
print("\nGenerating Embeddings for Training Data...")
X_train_embeddings = generate_embeddings(X_train.tolist(), tokenizer, model)
```



```
Generating Embeddings for Training Data...
```

```
print("Generating Embeddings for Testing Data...")
X_test_embeddings = generate_embeddings(X_test.tolist(), tokenizer, model)
```

✓ Data Loader Preparation

```
from torch.utils.data import DataLoader, Dataset
```

```
class SentimentDataset(Dataset):
    def __init__(self, embeddings, labels):
        self.embeddings = torch.tensor(embeddings, dtype=torch.float32)
        # Convert labels to PyTorch long tensors
        self.labels = torch.tensor(labels.values, dtype=torch.long)
```

```
    def __len__(self):
        return len(self.labels)
```

```
    def __getitem__(self, idx):
        return self.embeddings[idx], self.labels[idx]
```

```
# Create Dataset
```

```
train_dataset = SentimentDataset(train_datasetX_train_embeddings, batch_size=32
test_loader = DataLoader(test_dataset, batch_size=32)
, y_train)
test_dataset = SentimentDataset(X_test_embeddings, y_test)
```

```
# Create DataLoaders
```

```
train_loader = DataLoader(
```

✓ Sentiment Classifier

```

class SentimentClassifier(nn.Module):
    def __init__(self, input_dim, num_classes):
        super(SentimentClassifier, self).__init__()
        self.fc = nn.Sequential(
            nn.Linear(input_dim, 128),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(128, num_classes)
        )

    def forward(self, x):
        return self.fc(x)

input_dim = X_train_embeddings.shape[1]
num_classes = 3 # Negative, Neutral, Positive
model = SentimentClassifier(input_dim, num_classes)

criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)

```

✓ Model Training

```

print("\nTraining the Model...")
epochs = 100
for epoch in range(epochs):
    model.train()
    epoch_loss = 0
    correct = 0
    total = 0

    for embeddings, labels in train_loader:
        optimizer.zero_grad()
        outputs = model(embeddings)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        epoch_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

    print(f"Epoch {epoch + 1}/{epochs}, Loss: {epoch_loss:.4f}, Accuracy: {corr

```

✓ Model Evaluation

```
print("\nEvaluating the Model...")
model.eval()
all_preds = []
all_labels = []
with torch.no_grad():
    for embeddings, labels in test_loader:
        outputs = model(embeddings)
        _, predicted = torch.max(outputs, 1)
        all_preds.extend(predicted.tolist())
        all_labels.extend(labels.tolist())

print("\nClassification Report:")
print(classification_report(all_labels, all_preds, target_names=["Negative", "N
```



Evaluating the Model...


Classification Report:

	precision	recall	f1-score	support
Negative	0.71	0.66	0.68	1076
Neutral	0.40	0.20	0.27	800
Positive	0.92	0.97	0.95	8124
accuracy			0.88	10000
macro avg	0.68	0.61	0.63	10000
weighted avg	0.86	0.88	0.86	10000

```

conf_matrix = confusion_matrix(all_labels, all_preds)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=["Negat
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```




Start coding or [generate](#) with AI.

Visualizations


```
# Generate Word Cloud
all_words = ' '.join(data['text_cleaned'])
wordcloud = WordCloud(width=800, height=400, background_color='white', colormap=

# Plot Word Cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title("Most Frequent Words in Reviews", fontsize=16)
plt.show()
```



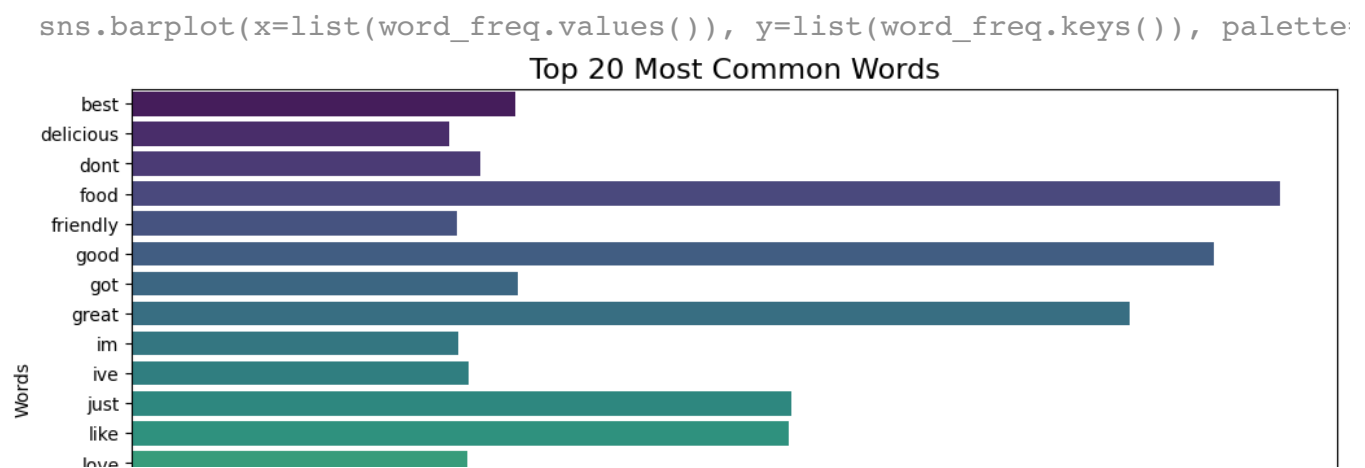
```
from sklearn.feature_extraction.text import CountVectorizer
```

```
vectorizer = CountVectorizer(stop_words='english', max_features=20)
word_counts = vectorizer.fit_transform(data['text_cleaned'])
word_freq = dict(zip(vectorizer.get_feature_names_out(), word_counts.toarray().
```

```
# Bar Plot
```

```
plt.figure(figsize=(12, 6))
sns.barplot(x=list(word_freq.values()), y=list(word_freq.keys()), palette='viri
plt.title("Top 20 Most Common Words", fontsize=16)
plt.xlabel("Frequency")
plt.ylabel("Words")
plt.show()
```

```
↪ /tmp/ipykernel_1703355/357633920.py:10: FutureWarning:
    Passing `palette` without assigning `hue` is deprecated and will be removed
```



```
# Sentiment Distribution
```

```
sentiment_counts = data['sentiment_label'].value_counts()
```

```
# Pie Chart
```


```
plt.figure(figsize=(8, 6))
plt.pie(sentiment_counts, labels=["Positive", "Neutral", "Negative"], autopct='
plt.title("Sentiment Distribution")
plt.show()
```

```
# Bar Plot
```

```
plt.figure(figsize=(8, 6))
sns.barplot(x=sentiment_counts.index, y=sentiment_counts.values, palette="virid
plt.xticks([0, 1, 2], ["Positive", "Neutral", "Negative"])
plt.title("Sentiment Distribution")
plt.xlabel("Sentiment")
plt.ylabel("Count")
plt.show()
```



Sentiment Distribution



/tmp/ipykernel_1703355/969110400.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed

```
sns.barplot(x=sentiment_counts.index, y=sentiment_counts.values, palette=
```


Sentiment Distribution



```
data['date'] = pd.to_datetime(data['date']) # Ensure date is in datetime format
data['month_year'] = data['date'].dt.to_period('M') # Group by month and year
sentiment_trend = data.groupby('month_year')['sentiment_label'].mean().reset_index
```


```
# Line Plot
```

```
plt.figure(figsize=(12, 6))
plt.plot(sentiment_trend['month_year'].astype(str), sentiment_trend['sentiment_label'])
plt.title("Average Sentiment Over Time", fontsize=16)
plt.xlabel("Date")
plt.ylabel("Average Sentiment (0=Negative, 1=Neutral, 2=Positive)")
plt.xticks(rotation=90)
plt.grid()
plt.show()
```



```
# Correlation Heatmap
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(data[['stars', 'useful', 'funny', 'cool']].corr(), annot=True, cmap=
plt.title("Correlation Between Numeric Features", fontsize=16)
plt.show()
```



```
data['review_length'] = data['text_cleaned'].apply(len)
```

```
# Box Plot
```

```
plt.figure(figsize=(10, 6))
```

```
sns.boxplot(x='sentiment_label', y='review_length', data=data, palette="viridis
```

```
plt.title("Review Length by Sentiment", fontsize=16)
```

```
plt.xlabel("Sentiment")
```

```
plt.ylabel("Review Length")
```


```
plt.xticks([0, 1, 2], ["Negative", "Neutral", "Positive"])
```

```
plt.show()
```

➞ /tmp/ipykernel_1703355/835442337.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed

```
sns.boxplot(x='sentiment_label', y='review_length', data=data, palette="v
```




```
# Group by Stars
sentiment_by_stars = data.groupby('stars')['sentiment_label'].mean()

# Bar Plot
plt.figure(figsize=(10, 6))
sns.barplot(x=sentiment_by_stars.index, y=sentiment_by_stars.values, palette="v")
plt.title("Average Sentiment by Star Ratings", fontsize=16)
plt.xlabel("Stars")
plt.ylabel("Average Sentiment (0=Negative, 1=Neutral, 2=Positive)")
plt.show()
```

➞ /tmp/ipykernel_1703355/1853316778.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed

```
sns.barplot(x=sentiment_by_stars.index, y=sentiment_by_stars.values, palette="v")
```



```

from wordcloud import WordCloud

# Filter Text Based on Sentiments
positive_text = ' '.join(data[data['sentiment_label'] == 2]['text_cleaned'])
neutral_text = ' '.join(data[data['sentiment_label'] == 1]['text_cleaned'])
negative_text = ' '.join(data[data['sentiment_label'] == 0]['text_cleaned'])

# Generate Word Clouds
positive_wc = WordCloud(width=800, height=400, background_color='white', color=
neutral_wc = WordCloud(width=800, height=400, background_color='white', color=
negative_wc = WordCloud(width=800, height=400, background_color='white', color=

# Plot Word Clouds
plt.figure(figsize=(16, 8))

plt.subplot(1, 3, 1)
plt.imshow(positive_wc, interpolation='bilinear')
plt.title("Positive Sentiment Word Cloud", fontsize=16)
plt.axis('off')

plt.subplot(1, 3, 2)
plt.imshow(neutral_wc, interpolation='bilinear')
plt.title("Neutral Sentiment Word Cloud", fontsize=16)
plt.axis('off')

plt.subplot(1, 3, 3)
plt.imshow(negative_wc, interpolation='bilinear')
plt.title("Negative Sentiment Word Cloud", fontsize=16)
plt.axis('off')

plt.tight_layout()
plt.show()

```



```

from sklearn.feature_extraction.text import CountVectorizer

# Function to Get Most Frequent Words
def get_most_frequent_words(texts, top_n=10):
    vectorizer = CountVectorizer(stop_words='english', max_features=top_n)
    word_counts = vectorizer.fit_transform(texts)
    word_freq = dict(zip(vectorizer.get_feature_names_out(), word_counts.toarray()))
    return word_freq

# Get Most Frequent Words
positive_words = get_most_frequent_words(data[data['sentiment_label'] == 2]['te

```

```

neutral_words = get_most_frequent_words(data[data['sentiment_label'] == 1]['text'])
negative_words = get_most_frequent_words(data[data['sentiment_label'] == 0]['text'])

# Plot Most Frequent Words
plt.figure(figsize=(18, 6))

plt.subplot(1, 3, 1)
sns.barplot(x=list(positive_words.values()), y=list(positive_words.keys()), palette='magma')
plt.title("Most Frequent Words (Positive Sentiment)", fontsize=16)
plt.xlabel("Frequency")
plt.ylabel("Words")

plt.subplot(1, 3, 2)
sns.barplot(x=list(neutral_words.values()), y=list(neutral_words.keys()), palette='magma')
plt.title("Most Frequent Words (Neutral Sentiment)", fontsize=16)
plt.xlabel("Frequency")
plt.ylabel("Words")

plt.subplot(1, 3, 3)
sns.barplot(x=list(negative_words.values()), y=list(negative_words.keys()), palette='magma')
plt.title("Most Frequent Words (Negative Sentiment)", fontsize=16)
plt.xlabel("Frequency")
plt.ylabel("Words")

plt.tight_layout()
plt.show()

```

⚠ /tmp/ipykernel_1703355/1829494720.py:19: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed

```

sns.barplot(x=list(positive_words.values()), y=list(positive_words.keys()), palette='magma')
/tmp/ipykernel_1703355/1829494720.py:25: FutureWarning:

```


Passing `palette` without assigning `hue` is deprecated and will be removed

```

sns.barplot(x=list(neutral_words.values()), y=list(neutral_words.keys()), palette='magma')
/tmp/ipykernel_1703355/1829494720.py:31: FutureWarning:

```

Passing `palette` without assigning `hue` is deprecated and will be removed



Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.

Start coding or generate with AI.